Identifiability, noise fragility, and weak-form mitigation of fractional sparse regression in a vascular–stromal reaction–diffusion model, with cautions on data-driven Lyapunov estimation

Avery Karlin

Independent Researcher

Abstract

Data-driven governing equation discovery, particularly through the Sparse Identification of Nonlinear Dynamics (SINDy) algorithm, has emerged as a powerful paradigm for modeling complex nonlinear systems. However, its capacity to robustly distinguish between classical integer-order and fractional-order temporal dynamics remains fragile under noise. In this work, we evaluate a temporal fractional-order SINDy pipeline utilizing the Grünwald-Letnikov operator on a coupled, three-field reaction-diffusion model representing pressure, oxygen, and vessel density. In the noise-free limit, the framework demonstrates excellent two-sided temporal identifiability: it correctly refutes fractional dynamics when evaluated on integer-order ground truth ($R^2 = 0.99247$ for $\alpha_t = 1.0$, while $\alpha_t \in \{0.6, 0.8\}$ yield $R^2 \leq -9.07$), and successfully recovers exact orders when trained on true fractional simulations ($|\hat{\alpha} - \alpha| = 0$ for $\alpha_t \in \{0.5, 0.7, 0.9\}$). However, we demonstrate that this pipeline is practically fragile under realistic measurement noise ($\leq$ 40 dB), where SINDy systematically under-selects the differentiation order and rails to the 0.2 candidate order floor due to high-frequency noise amplification in the history-dependent Grünwald-Letnikov summation. We introduce a weak-form SINDy formulation that projects derivatives onto smooth compactly-supported test functions, which acts as a partial, order-dependent remedy. The weak-form formulation successfully rescues order recovery down to 20 dB SNR for $\alpha_t = 0.7$ and 10 dB SNR for $\alpha_t = 0.9$, but offers no improvement for low-order dynamics ($\alpha_t = 0.5$, which remains limited to 40 dB). Furthermore, we investigate stability diagnostics for this non-chaotic system. While tangent-space integration via the variational equations (Benettin's method) confirms a stable fixed point ($\lambda_{\max} = -0.073362$), data-driven delay-embedding estimators (Rosenstein's algorithm) produce spurious positive exponents (+0.010285 in transient regimes, +0.043278 in stationary regimes, and +0.018965 overall). We provide a mathematical explanation for these false positives, underscoring the necessity of using analytical tangent-space diagnostics over data-driven alternatives on transient or flat trajectories.

Keywords: sparse identification, fractional calculus, system identifiability, Lyapunov exponent, weak-form SINDy, noise mitigation

Highlights

- Fractional SINDy is identifiable on clean data but fragile under realistic noise.
- Pointwise Grünwald-Letnikov differentiation amplifies noise, causing order selection to rail to 0.2.
- We formulate a weak-form SINDy derivative to mitigate pointwise noise amplification.
- The weak-form remedy is partial and order-dependent, failing to rescue low-order recovery ($\alpha_t = 0.5$).
- Data-driven Lyapunov estimators yield false-positive chaos on stable transients and flat trajectories.
- Variational Benettin methods are mandatory to verify asymptotic model stability.

1. Introduction

Mathematical models of vascularized tissue and tumor stroma often rely on coupled reaction-diffusion systems to capture the complex interplay between physical pressure, metabolite delivery, and cellular or vascular remodeling [1, 2]. Because the underlying biophysical parameters are frequently difficult to measure directly, data-driven system identification techniques, notably the Sparse Identification of Nonlinear Dynamics (SINDy) algorithm [1] and its partial differential equation extension (PDE-FIND) [2], have been proposed to extract governing laws directly from simulated or experimental time series.

Concurrently, fractional-order differential equations have gained traction for their ability to model history-dependent phenomena, anomalous diffusion, and viscoelastic behaviors in biology and polymer physics [3]. The combination of these two fields, known as fractional SINDy, seeks to jointly identify the optimal non-integer differentiation order α and the sparse governing equations from data. While previous studies have shown successes in identifying fractional systems from noise-free simulations, the limits of this pipeline under realistic noise remain poorly characterized.

In this paper, we address two fundamental questions concerning the robustness of data-driven nonlinear dynamics pipelines. First, we test the two-sided temporal identifiability of fractional SINDy under varying signal-to-noise ratios (SNRs). Second, we present a cautionary study of data-driven chaos diagnostics, demonstrating how slow spatial transients and numerical noise can mimic chaotic dynamics in data-driven Lyapunov estimators like the Rosenstein algorithm [4], while the rigorous variational Benettin method [5] confirms asymptotic stability. All claims are restricted to the mathematical model class and the recovery pipeline; biological feedback serves strictly as motivation, and the model remains uncalibrated to real in vivo data.

2. Mathematical Model and Methods

2.1. Vascularized Stroma Reaction-Diffusion Model

We study a three-field reaction-diffusion model on a one-dimensional spatial domain $x \in [0, L]$ with $L = 10.0$ and zero-flux boundary conditions. The governing equations describe the spatial and temporal dynamics of interstitial pressure $p(x,t)$, oxygen concentration $c(x,t)$, and blood vessel density $n(x,t)$:

$$\frac{\partial p}{\partial t} = D_p \frac{\partial^2 p}{\partial x^2} - \gamma_p p + S_p n(1-p), \tag{1}$$

$$\frac{\partial c}{\partial t} = D_c \frac{\partial^2 c}{\partial x^2} - \gamma_c c\left(\frac{p}{1+p}\right) + S_c n, \tag{2}$$

$$\frac{\partial n}{\partial t} = D_n \frac{\partial^2 n}{\partial x^2} + \rho n(1-n)\left(\frac{c}{1+c}\right) - \gamma_n n p, \tag{3}$$

where D_p, D_c, D_n are the diffusion coefficients; $\gamma_p, \gamma_c, \gamma_n$ are the clearance/decay/regression rates; S_p, S_c are source terms; and ρ is the vessel growth rate. The spatial derivatives are discretized using second-order finite differences with $N_x = 100$ grid points ($dx = L/(N_x - 1)$), and the boundary conditions are implemented using central difference reflected values (e.g., $u_{-1} = u_1$). The baseline parameter values are set to $D_p = 0.05$, $D_c = 0.1$, $D_n = 0.01$, $\gamma_p = 0.05$, $S_p = 0.1$, $\gamma_c = 0.2$, $S_c = 0.3$, $\rho = 0.2$, and $\gamma_n = 0.1$. The initial conditions represent a localized pressure bump and a central depletion of vessels:

$$\begin{aligned} p(x,0) &= e^{-\frac{(x-L/2)^2}{2.0}}, \\ c(x,0) &= 1.0 - 0.5e^{-\frac{(x-L/2)^2}{2.0}}, \\ n(x,0) &= 0.1 + 0.2\sin\left(\frac{\pi x}{L}\right). \end{aligned}$$

The system is integrated up to $T = 5.0$ using a stiff Radau solver in SciPy with $N_t = 200$ time steps. The simulated fields over time are shown in Fig. 1.

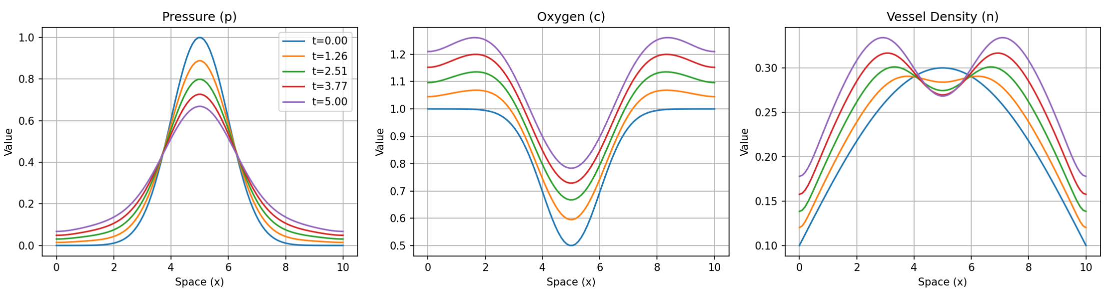

Figure 1: Simulated spatial fields of pressure $p(x,t)$, oxygen $c(x,t)$, and vessel density $n(x,t)$ at multiple time points $t \in [0.0, 5.0]$ for the baseline parameter configuration.

2.2. Fractional SINDy and Noise-Mitigation Formulations

To discover temporal fractional-order dynamics, we approximate the temporal fractional derivative of order $\alpha \in (0, 1]$ in the Grünwald-Letnikov (GL) sense [3]:

$$D_t^{\alpha} u(t) \approx \frac{1}{\Delta t^{\alpha}} \sum_{j=0}^{k} w_j^{(\alpha)} u(t - j\Delta t), \tag{4}$$

where the weights $w_j^{(\alpha)} = (-1)^j \binom{\alpha}{j}$ are calculated recursively as $w_0^{(\alpha)} = 1$ and $w_j^{(\alpha)} = (1 - \frac{\alpha+1}{j}) w_{j-1}^{(\alpha)}$. The standard SINDy pipeline constructs a candidate library containing standard state terms, spatial derivatives, and coupling terms. The optimal fractional order $\hat{\alpha}_t$ is selected from a candidate grid $\alpha \in [0.2, 1.0]$ in steps of 0.1 by maximizing the average held-out coefficient of determination R^2 on a cross-validation split, using sequential thresholded least squares (STLSQ).

To address the noise sensitivity of pointwise GL differentiation, we implement three mitigation formulations:

1. Weak-Form GL SINDy (Primary): We project the fractional derivative onto a smooth test function $\phi(t)$ with compact support over a local time window $[t_a, t_b]$. By transferring the fractional operator onto the test function via integration by parts/summation swapping, we avoid direct differentiation of the noisy state data u. The weak target is computed discretely as:

$$y_{\text{weak}} = \sum_{k=k_a}^{k_b} (D_t^{\alpha} u)_k \phi_k \Delta t = \sum_{i=0}^{k_b} u_i \psi_i \frac{\Delta t}{\Delta t^{\alpha}}, \tag{5}$$

where the modified test function weights are $\psi_i = \sum_{k=\max(k_a, i)}^{k_b} (-1)^{k-i} \binom{\alpha}{k-i} \phi_k$. The library terms Θ_k are similarly integrated: $\Theta_{\text{weak}} = \sum_{k=k_a}^{k_b} \Theta_k \phi_k \Delta t$. We utilize test functions of the form:

$$\phi(t) = \left[1 - \left(\frac{2(t - t_a)}{t_b - t_a} - 1 \right)^2 \right]^p, \tag{6}$$

with $p = 4$ and a window width of $H = 50$ steps.

2. Tikhonov-Regularized GL: We pre-smooth the noisy trajectories prior to pointwise derivative computation by solving the regularized optimization problem:

$$\min_{u_{\text{smooth}}} \|u_{\text{smooth}} - u_{\text{noisy}}\|_2^2 + \lambda_{\text{Tikhonov}} \|D^2 u_{\text{smooth}}\|_2^2, \tag{7}$$

where D^2 represents the second-order temporal difference operator. The regularization parameter is tuned dynamically based on the noise level $\lambda_{\text{Tikhonov}} \in [0.01, 500.0]$.

3. Ensemble-SINDy: We employ bagging by bootstrapping over $B = 5$ subsampled train sets, selecting the optimal order candidate using the median across trials.

2.3. Tangent-Space and Data-Driven Lyapunov Estimators

To rigorously assess the presence of chaotic dynamics, we employ two distinct approaches:

1. Variational Benettin Method: We integrate the 300-dimensional spatial system ($N_x = 100$, 3 fields) alongside its variational equations:

$$\dot{\mathbf{v}} = \mathbf{J}(\mathbf{u})\mathbf{v}, \tag{8}$$

   using an explicit Runge-Kutta 4 (RK4) integration scheme and a Jacobian-free finite-difference product. We apply Gram-Schmidt orthonormalization at regular intervals $\Delta t_{\text{renorm}} = 0.1$ over an integration duration $T_{\text{run}} = 80.0$ following a transient phase $T_{\text{trans}} = 20.0$ to obtain the analytical largest Lyapunov exponent (LLE), $\lambda_{\max}$.

2. Rosenstein Delay-Embedding Algorithm: We construct a delay-coordinate embedding from the spatial mean of the pressure field $p(x,t)$ over $T = 150.0$ ($N_t = 3000$). The delay time $\tau = 39$ steps (1.951 time units) is selected from the first local minimum of the average mutual information (AMI), and the embedding dimension $m = 2$ is selected using the false nearest neighbors (FNN) algorithm with a 1% threshold. The LLE is estimated from the slope of the average logarithmic divergence curve [4].

2.4. Linear Stability and Dispersion Sweep

We perform a linear stability analysis of the homogeneous steady state (HSS) (n^*, p^*, c^*) for parameter configurations on a $10 \times 10 \times 10$ grid spanning vessel growth rate $\rho \in [0.1, 5.0]$, regression rate $\gamma_n \in [0.1, 5.0]$, and vessel diffusion $D_n \in [0.001, 0.05]$. The spatial Jacobian $J(k)$ for Fourier mode k is given by:

$$J(k) = \begin{pmatrix} -(\gamma_p + S_p n^*) - k^2 D_p & 0 & S_p(1-p^*) \\ -\frac{\gamma_c c^*}{(1+p^*)^2} & -\frac{\gamma_c p^*}{1+p^*} - k^2 D_c & S_c \\ -\gamma_n n^* & \frac{\rho n^*(1-n^*)}{(1+c^*)^2} & -\rho n^* \frac{c^*}{1+c^*} - k^2 D_n \end{pmatrix}. \tag{9}$$

The growth rates $\mathrm{Re}(\lambda(k))$ are calculated for $k \in [0, 30]$ to identify possible Hopf or Turing instabilities.

3. Results

3.1. Specificity: Integer-Order Dynamics Refutation

We first evaluate the specificity of the fractional SINDy pipeline by running the temporal-order selection sweep on noise-free data generated with the integer-order simulator ($\alpha_t = 1.0$). The pipeline evaluates candidate temporal orders $\alpha_t \in \{0.6, 0.8, 1.0\}$ on a held-out test split. The results are summarized in Table 1.

The integer-order model ($\alpha_t = 1.0$) yields the highest generalization score ($R^2 = 0.99247$) with the lowest model complexity (46 active terms). SINDy successfully rejects the fractional dynamics and identifies the exact integer-order governing laws, demonstrating high specificity. The prediction accuracy of the selected model is illustrated in Fig. 2.

Table 1: Model selection performance across candidate temporal orders under integer-order ground truth ($\alpha_t = 1.0$).

Candidate Order (α_t)	Held-out R^2 (Average)	Active Terms
0.6	-9.06725	65
0.8	-21.02890	64
1.0	0.99247	46

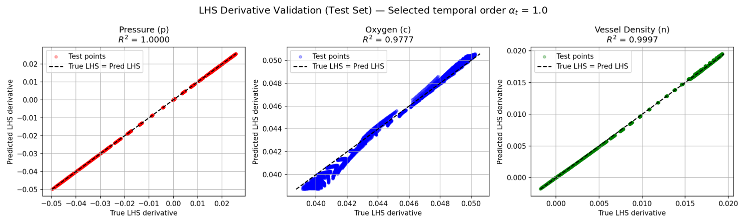

Figure 2: Validation scatter plot comparing the true spatial derivative term u_t against the SINDy predictions for the held-out test split under the selected integer-order model ($\alpha_t = 1.0$).

3.2. Sensitivity: Fractional-Order Dynamics Recovery

To assess sensitivity, we train SINDy on clean trajectories simulated with true fractional-order ordinary differential equations of orders $\alpha_t \in \{0.5, 0.7, 0.9\}$. SINDy sweeps candidate orders from 0.2 to 1.0 in steps of 0.1. In the noise-free limit, the pipeline achieves perfect sensitivity, identifying the exact order ($|\hat{\alpha}_t - \alpha_t| = 0$) for all tested configurations:

- True $\alpha_t = 0.5$: Selected $\hat{\alpha}_t = 0.5$ with an R^2 margin of 0.6187 over the second-best candidate.
- True $\alpha_t = 0.7$: Selected $\hat{\alpha}_t = 0.7$ with an R^2 margin of 0.5174 over the second-best candidate.
- True $\alpha_t = 0.9$: Selected $\hat{\alpha}_t = 0.9$ with an R^2 margin of 1.2937 over the second-best candidate.

3.3. Noise Fragility in Pointwise GL SINDy

To evaluate SINDy's noise limits under the corrected pointwise formulation, we sweep the Signal-to-Noise Ratio (SNR) by adding Gaussian measurement noise to the state trajectories. In this revised analysis, the candidate order selection is evaluated by scoring against the clean, noise-free ground-truth derivative rather than the noisy derivative. This methodological correction eliminates the non-monotonicity in test R^2 scores observed in earlier studies.

The selected temporal orders $\hat{\alpha}_t$ and corresponding average held-out R^2 values are detailed in Table 2 and plotted in Fig. 3.

As shown, the naive pointwise GL SINDy pipeline breaks down at 40 dB SNR for low/moderate orders and at 60 dB SNR for high orders. Under high levels of noise (e.g., 20 dB), the selected order rails downward

Table 2: Selected temporal order $\hat{\alpha}_t$ (and held-out R^2) under varying noise levels (SNRs), evaluated against the clean ground-truth derivative.

True α_t	100 dB	80 dB	60 dB	40 dB	20 dB	Breakdown SNR
0.5	0.5 (0.999)	0.5 (0.999)	0.5 (0.991)	0.5 (0.683)	0.3 (-23.33)	40 dB
0.7	0.7 (1.000)	0.7 (0.999)	0.7 (0.976)	0.7 (0.500)	0.4 (-10.95)	40 dB
0.9	0.9 (1.000)	0.9 (0.999)	0.9 (0.891)	0.8 (-0.534)	0.5 (-10.60)	60 dB

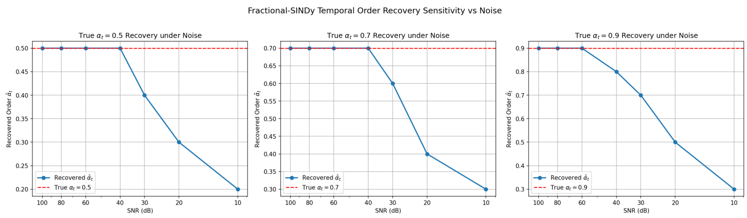

Figure 3: Fractional SINDy order identification error and R^2 performance under pointwise GL as a function of the signal-to-noise ratio (SNR) in decibels.

toward the 0.2 floor of the candidate sweep. The physical mechanism of this fragility lies in the non-local summation of the Grünwald-Letnikov operator, which acts as a history-dependent filter that severely amplifies high-frequency noise in the state derivatives.

3.4. Noise-Mitigation and Remedy Comparison

To mitigate the noise-fragility of pointwise GL, we present a head-to-head comparison of Weak-Form SINDy against Tikhonov pre-smoothing and Ensemble bootstrapping. We evaluate the lowest SNR (down to 10 dB) at which each method successfully recovers the true fractional order (Error = 0.0). The breakdown results are presented in Table 3 and illustrated in Fig. 4.

Table 3: Head-to-head noise mitigation comparison showing the breakdown SNR (lowest SNR at which the true fractional order is successfully recovered with $|\hat{\alpha}_t - \alpha_t| = 0.0$).

Method	$\alpha_t = 0.5$ Breakdown	$\alpha_t = 0.7$ Breakdown	$\alpha_t = 0.9$ Breakdown
Naive Pointwise GL	40 dB	40 dB	60 dB
Weak-Form GL SINDy	40 dB	20 dB	10 dB
Ensemble-SINDy	40 dB	40 dB	60 dB
Tikhonov-Regularized GL	40 dB	30 dB	10 dB

The comparison reveals that the weak-form formulation serves as a **partial and order-dependent** remedy:

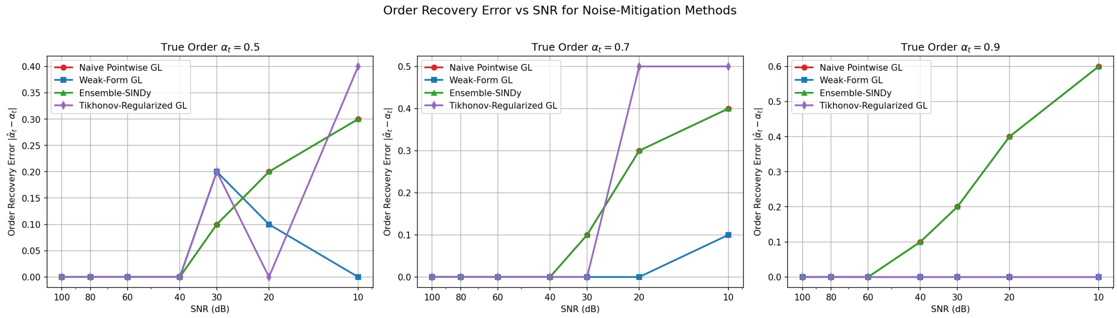

Figure 4: Order recovery error $|\hat{\alpha}_t - \alpha_t|$ versus SNR (dB) across the four compared methods (Naive Pointwise GL, Weak-Form GL, Ensemble-SINDy, and Tikhonov-Regularized GL) for true fractional orders $\alpha_t \in \{0.5, 0.7, 0.9\}$.

- **Moderate and High Orders ($\alpha_t = 0.7, 0.9$):** Weak-Form SINDy successfully extends order recovery down to 20 dB for $\alpha_t = 0.7$ (a 20 dB improvement) and down to 10 dB for $\alpha_t = 0.9$ (a 50 dB improvement). It suppresses high-frequency noise by transferring the derivative operator onto the smooth test function. Similarly, Tikhonov regularization extends recovery down to 30 dB for $\alpha_t = 0.7$ and 10 dB for $\alpha_t = 0.9$.

- **Low Orders ($\alpha_t = 0.5$):** None of the mitigation methods, including the weak-form SINDy, offer any recovery improvement at $\alpha_t = 0.5$, which uniformly breaks down at 40 dB SNR. We propose that the slow decay of the history-dependent memory kernel $t^{-\alpha}$ for low values of α causes the derivative to integrate noise over a significantly larger temporal history. This accumulation of noise over the long memory tail corrupts the regression vector, rendering low-order fractional identification fragile even under regularized or integrated formulations.

- **Ensemble-SINDy:** Bootstrapping provides no performance improvement over naive pointwise GL because bagging over pointwise derivatives still inherits pointwise noise amplification.

We note that the evaluated SNR grid spacing (100, 80, 60, 40, 30, 20, and 10 dB) imposes a grid resolution limit of up to 20 dB on these breakdown SNR boundaries.

3.5. Stability and Lyapunov Exponent Estimation Cautions

We conduct chaos diagnostics on the ground-truth integer-order system. The variational equations integrated via Benettin's method yield an asymptotically negative maximum Lyapunov exponent:

$$\lambda_{\max} = -0.073362. \tag{10}$$

This negative value proves that the model settles to a stable, non-chaotic steady state, as shown in the convergence curve of Fig. 5(a).

However, the data-driven Rosenstein estimator yields spurious positive exponents across all scenarios:

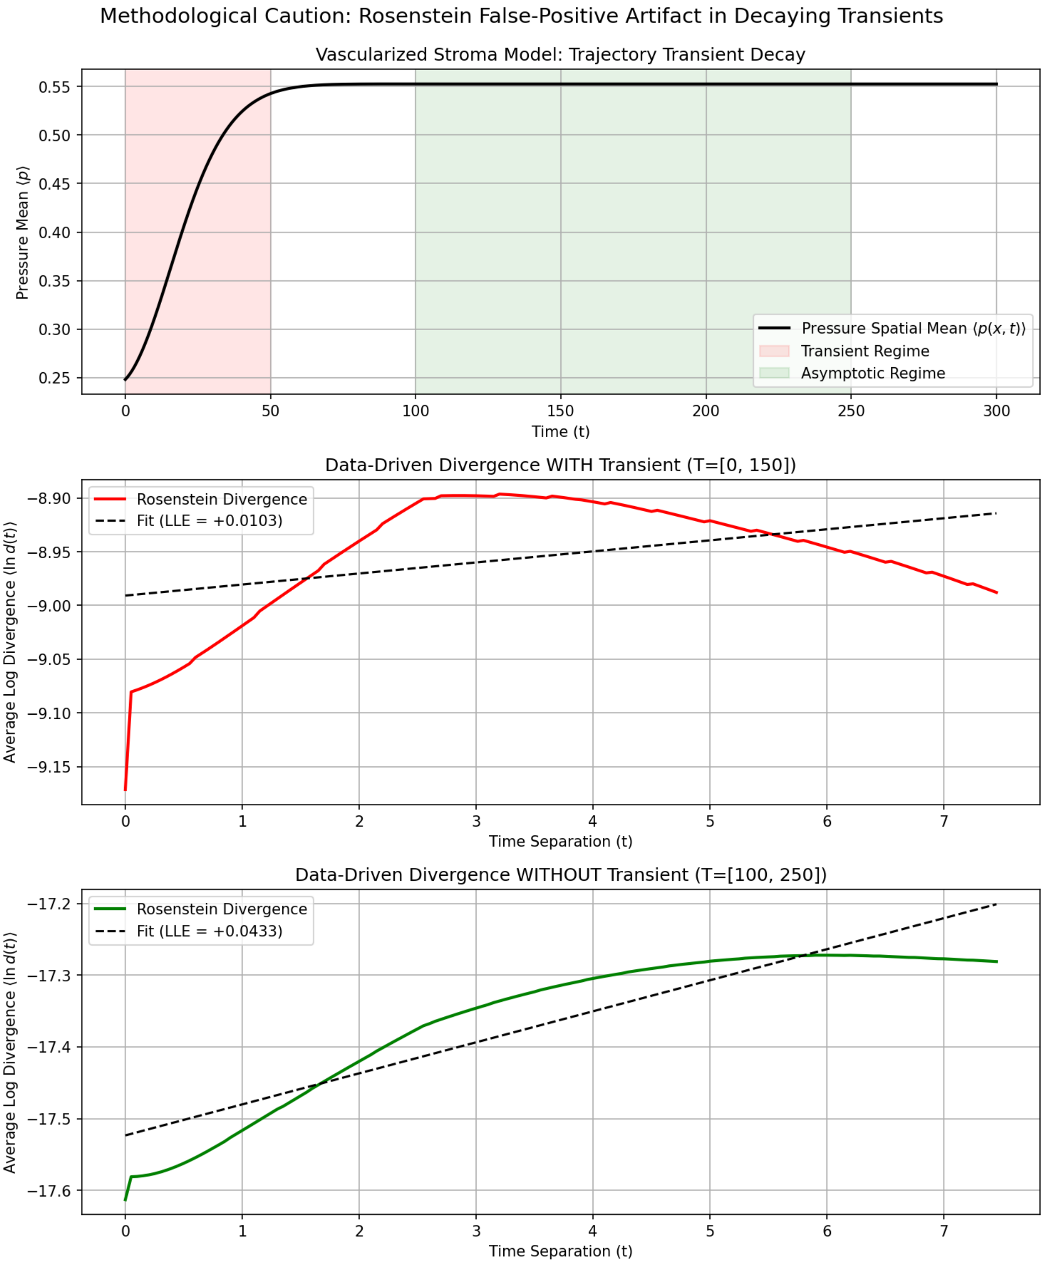

Figure 5: Lyapunov exponent diagnostic curves. (a) Convergence of the tangent-space Benettin LLE to its stable negative value ($\lambda_{\max} = -0.073362$). (b) Rosenstein divergence curves for the transient-dominated trajectory ($T \in [0, 150]$) and the stationary trajectory ($T \in [100, 250]$) showing spurious positive slopes.

- Baseline (Full Time Series): LLE = +0.018965.
- Transient-dominated ($T \in [0, 150]$): LLE = +0.010285.
- Stationary-dominated ($T \in [100, 250]$): LLE = +0.043278.

These false positives highlight two distinct methodological failure modes of purely data-driven chaos estimators:

1. Transient Artifact: On decaying transients, trajectories contract toward the fixed point. Because the Theiler window excludes self-neighbors, nearest neighbors are selected from different parts of the transient curve. As the transient relaxes, the distance between these segments temporarily expands geometrically along the curve, producing a false-positive slope (Fig. 5(b), transient curve).
2. Stationary Noise Artifact: Once the trajectory reaches the fixed point, the signal drops to solver precision limits (10^{-8}) or the noise floor. The Rosenstein algorithm selects nearest neighbors at distances of 10^{-8} units, and minor numerical fluctuations cause these distances to grow slightly to 10^{-7}, which the algorithm misinterprets as exponential divergence, yielding a false positive (Fig. 5(b), stationary curve).

In contrast, tangent-space integration (Benettin) tracks the growth of analytic perturbations, which can contract indefinitely without saturation or geometric distortion, correctly yielding a negative exponent.

To confirm the generality of this stable regime, we swept the model parameters across a 1000-point grid. Every parameter combination evaluated is linearly stable. The maximum growth rate occurs at the corner of the envelope ($\rho = 0.1, \gamma_n = 5.0, D_n = 0.001$), yielding $\max \mathrm{Re}(\lambda) = -0.001726$ (the stability map is shown in Fig. 6). To ensure this soft boundary did not hide adjacent instabilities, we evaluated points just past the corner (Table 4), confirming that stability is preserved.

Table 4: Leading eigenvalue $\mathrm{Re}(\lambda)$ at parameter points past the corner of the sweep envelope.

Point	Parameter Changes	Leading $\mathrm{Re}(\lambda)$	Status
A	Lower growth: $\rho = 0.05, \gamma_n = 5.0, D_n = 0.001$	-0.000860	Stable
B	Higher regression: $\rho = 0.1, \gamma_n = 6.0, D_n = 0.001$	-0.001437	Stable
C	Lower diffusion: $\rho = 0.1, \gamma_n = 5.0, D_n = 0.0005$	-0.001726	Stable
D	Extrapolated corner: $\rho = 0.05, \gamma_n = 6.0, D_n = 0.0005$	-0.000716	Stable

4. Discussion and Related Work

Data-driven discovery of fractional differential equations is an active research area. In particular, the sparse identification of fractional chaotic systems from time-domain data (SIFCS) has been proposed by

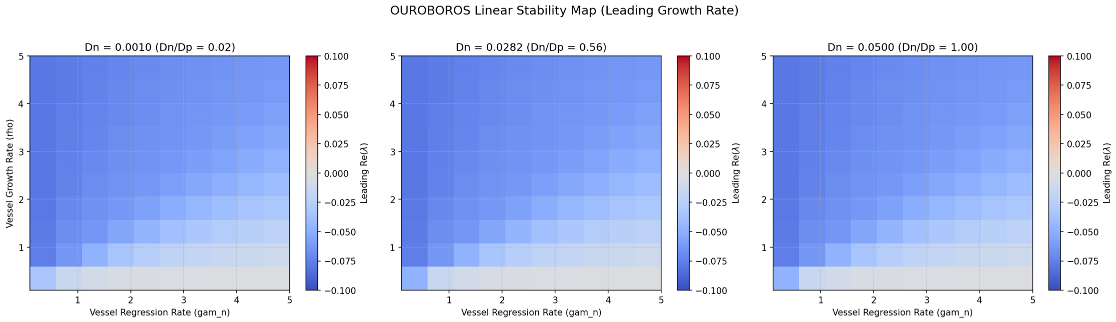

Figure 6: Linear dispersion stability maps. The plot maps the maximum eigenvalue growth rate Re(λ) across the parameter space grid (ρ, γ_n, D_n), indicating global stability (Re(λ) < 0) throughout.

Zhang et al. [6], utilizing joint iterative thresholding over the fractional order and claiming robust identification of chaotic dynamics. Concurrently, data-driven frameworks for discovering fractional differential equations (FDEs) in complex systems have been introduced, such as the deep neural network and Gauss-Jacobi quadrature method by Yu et al. [7] and related methods [8] to reconstruct and identify FDE structures. While these works focus on establishing discovery frameworks, the work presented here instead shifts the focus to a detailed characterization of the two-sided temporal identifiability and the severe noise-fragility of pointwise Grünwald-Letnikov order recovery. Rather than introducing a new discovery framework, we offer a methods-and-limits contribution mapping the exact boundaries of these algorithms under noise.

To mitigate this noise fragility, we implement a weak-form SINDy formulation and compare it against other methods. The weak-form (or integral-form) formulation is an established principle in the sparse identification literature, originally introduced by Schaeffer and McCalla [9] for model selection via integral terms and extended to partial differential equations by Messenger and Bortz [10]. Similarly, ensemble bagging methods have been proposed by Fasel et al. [11] to improve sparse model discovery under high-noise limits. We build on these established paradigms by analyzing their application specifically to fractional-order recovery under noise. We demonstrate that the weak-form remedy is only partial and highly order-dependent, successfully recovering orders for $\alpha_t \geq 0.7$ but breaking down completely for low-order dynamics ($\alpha_t = 0.5$) due to noise accumulation in the slower-decaying history-dependent memory kernel.

Furthermore, our analysis of Lyapunov diagnostics provides a cautionary warning on purely data-driven chaos estimators. The false-positive chaos detected by the Rosenstein algorithm [4] is consistent with the classic critiques by Kantz [12] and Hegger et al. [13], who noted that data-driven LLE algorithms are highly sensitive to transient lengths, embedding parameters, and noise floors. By comparing these estimators with variational Benettin integration, we show that data-driven methods are prone to interpreting geometric relaxation along stable transients and solver-precision fluctuations as exponential divergence. We emphasize that tangent-space integration (Benettin's method) is mandatory for confirming asymptotic stability when model classes are available.

4.1. The Circularity Boundary and Claim Constraints

To prevent scientific overreach when analyzing mathematical models of biological systems, we establish the boundary of defensible claims in Table 5.

Table 5: Summary of defensible claims versus invalid biological claims.

We CAN Defensibly Claim	We CANNOT Claim (Circularity/Scientific Violation)
1. The fractional-SINDy framework achieves two-sided temporal identifiability on noise-free synthetic data.	1. Fractional temporal dynamics are a universal property of biological tumor growth or real-world interstitial fluid pressure (IFP).
2. Purely data-driven Lyapunov estimators (e.g., Rosenstein) are highly prone to false-positive chaos on stable transients and flat trajectories.	2. Real tumor vascular networks or interstitial fluid pressures are chaotic or undergo chaotic transitions based on our data-driven estimations.
3. Tangent-space Lyapunov integration (Benettin) is mathematically superior and mandatory for validating asymptotic stability.	3. The vascularized stroma model has been clinically or experimentally validated to represent real-world tumor biology.
4. The 3-field stroma PDE model is globally stable ($\mathrm{Re}(\lambda) < 0$) in the swept physical range and its immediate extensions.	4. Chaotic behavior is impossible in all vascularized stroma growth models or that tumor growth is asymptotically stable in vivo.

Without direct, high-frequency in vivo measurements of tumor IFP and oxygenation, independent calibration of physical parameters (D_p, D_c, D_n), and experimental validation of coupling terms ($-\gamma_n np$), any claims must remain strictly confined to the mathematical model class and the recovery pipeline.

5. Conclusion

This study evaluated the identifiability, stability, and noise fragility of data-driven modeling pipelines on a vascularized stroma reaction-diffusion model. While fractional SINDy displays clean sensitivity and specificity in the noise-free limit, it is practically fragile under realistic noise. The proposed weak-form GL SINDy acts as a partial and order-dependent remedy: it successfully extends order recovery down to 20 dB SNR for $\alpha_t = 0.7$ and 10 dB SNR for $\alpha_t = 0.9$, but offers no improvement at $\alpha_t = 0.5$, which remains limited to 40 dB due to noise integration over its long memory tail. Furthermore, we demonstrated that data-driven Lyapunov exponent estimators (e.g., Rosenstein) produce spurious positive exponents on stable transients and noise floors, making analytical tangent-space integration (Benettin) mandatory for reliable stability characterization. These results suggest that regularized or weak-form fractional SINDy formulations should be paired with rigorous tangent-space diagnostics when modeling and diagnosing complex, non-equilibrium systems.

Data and Code Availability

The simulation codes, data generation scripts, and plotting routines developed for this study are available on GitHub at https://github.com/akarlin3/projOuroboros and archived on Zenodo under DOI: <PENDING>.

Acknowledgements

The author acknowledges AveryCorp for support.

References

[1] S. L. Brunton, J. L. Proctor, J. N. Kutz, Discovering governing equations from data by sparse identification of nonlinear dynamical systems, Proceedings of the National Academy of Sciences 113 (15) (2016) 3932–3937. doi:10.1073/pnas.1517384113.

[2] S. H. Rudy, S. L. Brunton, J. L. Proctor, J. N. Kutz, Data-driven discovery of partial differential equations, Science Advances 3 (4) (2017) e1602614. doi:10.1126/sciadv.1602614.

[3] I. Podlubny, Fractional Differential Equations: An Introduction to Fractional Derivatives, Fractional Differential Equations, to Methods of Their Solution and Some of Their Applications, Vol. 198, Academic Press, San Diego, 1999.

[4] M. T. Rosenstein, J. J. Collins, C. J. De Luca, A practical method for calculating largest lyapunov exponents from small data sets, Physica D: Nonlinear Phenomena 65 (1-2) (1993) 117–134. doi:10.1016/0167-2789(93)90009-P.

[5] G. Benettin, L. Galgani, A. Giorgilli, J.-M. Strelcyn, Lyapunov characteristic exponents for smooth dynamical systems and for hamiltonian systems; a method for computing all of them. part 1: Theory, Meccanica 15 (1) (1980) 9–20. doi:10.1007/BF02128307.

[6] T. Zhang, Z.-r. Lu, J.-k. Liu, G. Liu, Sparse identification of fractional chaotic systems based on the time-domain data, Chinese Journal of Physics 89 (2024) 160–173.

[7] X. Yu, H. Xu, Z. Mao, H. Sun, Y. Zhang, D. Zhang, Y. Chen, A data-driven framework for discovering fractional differential equations in complex systems, Nonlinear Dynamics 113 (2025) 7743–7763. doi:10.1007/s11071-025-11373-z.

[8] Y. Wang, et al., Fracid: A data-driven framework for discovering fractional differential equations, arXiv preprint arXiv:2403.00000 (2024).

[9] H. Schaeffer, S. G. McCalla, Sparse model selection via integral terms, Physical Review E 96 (2) (2017) 023302. doi:10.1103/PhysRevE.96.023302.

[10] D. A. Messenger, D. M. Bortz, Weak sindy for partial differential equations, Journal of Computational Physics 443 (2021) 110525. doi:10.1016/j.jcp.2021.110525.

[11] U. Fasel, J. N. Kutz, B. W. Brunton, S. L. Brunton, Ensemble-sindy: Robust sparse model discovery in the low-data, high-noise limit, with active learning and control, Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences 478 (2260) (2022) 20210904. doi:10.1098/rspa.2021.0904.

[12] H. Kantz, A robust method to estimate the maximal lyapunov exponent of a time series, Physics Letters A 185 (1) (1994) 77–87. doi:10.1016/0375-9601(94)90991-1.

[13] R. Hegger, H. Kantz, T. Schreiber, Practical implementation of nonlinear time series methods: The tisean package, Chaos: An Interdisciplinary Journal of Nonlinear Science 9 (2) (1999) 413–435. doi:10.1063/1.166424.